



VIET NAM NATIONAL UNIVERSITY HO CHI MINH CITY
INTERNATIONAL UNIVERSITY

Operating Systems

W1 - Introduction

Le Hai Duong, PhD. (lhduong@hcmiu.edu.vn) - 08/09/2021

Materials

- Textbook:
 - Arpaci-Dusseau, R. H., & Arpaci-Dusseau, A. C. (2018). *Operating systems: Three easy pieces*. Arpaci-Dusseau Books LLC. (Online book: <https://pages.cs.wisc.edu/~remzi/OSTEP/>)
- Other useful books:
 - Galvin, P. B., Gagne, G., & Silberschatz, A. (2018). *Operating system concepts* (10th ed.). John Wiley & Sons.
 - Tanenbaum, A. S., & Bos, H. (2015). *Modern operating systems* (4th ed.). Pearson.

Tentative Schedule

Session	Date	Title	Reading
		VIRTUALIZATION	
1	Sep 8, 2021	- Introduction	- Intro: pre , Ch1 , Ch2
2	Sep 15, 2021	- Processes	Ch3 , Ch4 , Ch5 , Ch6
3	Sep 22, 2021	- Schedule	Ch7 , Ch8 , Ch11
4	Sep 29, 2021	- Memory management - Paging	- Mem Mgmt: Ch13 , Ch15 , Ch16 - Paging: Ch18
5	Oct 6, 2021	- Paging: TLBs - Paging: Smaller Tables	- TLBs: Ch19 - Smaller: Ch20
		Beyond Physical: Swapping	Ch21 , Ch22
6	Oct 13, 2021	CONCURRENCY	
7	Oct 20, 2021	- Threads - Locks	- Threads: Ch25 , Ch26 - Locks: Ch28
8	Oct 27, 2021	- Locks, CVs - Condition variables	- Locks, CVs: Ch29 - Condition variables: Ch30
MIDTERM (1/11/2021 - 14/11/2021)			
9	Nov 17, 2021	Semaphores	Ch31
		Deadlocks	Ch32
10	Nov 24, 2021	PERSISTENCE	
11	Dec 1, 2021	- I/O + Disks - Disk schedule	- I/O: Ch36 - Disk, disk schedule: Ch37
12	Dec 8, 2021	- RAID - File systems	- RAID: Ch38 - File systems: Ch39
13	Dec 15, 2021	- File-System Implementation - Fast file system	- FS Impl: Ch40 - FFS: Ch41
14	Dec 22, 2021	Journaling	Ch42
15	Dec 29, 2021	- LFS - SSD	- LFS: Ch43 - SSD: Ch44
FINAL (03/01/2021 - 21/01/2021)			

Gradings

- Labs, homeworks, quizzes 30%
- Midterm exam 30%
- Final exam 40%

Computer Organization

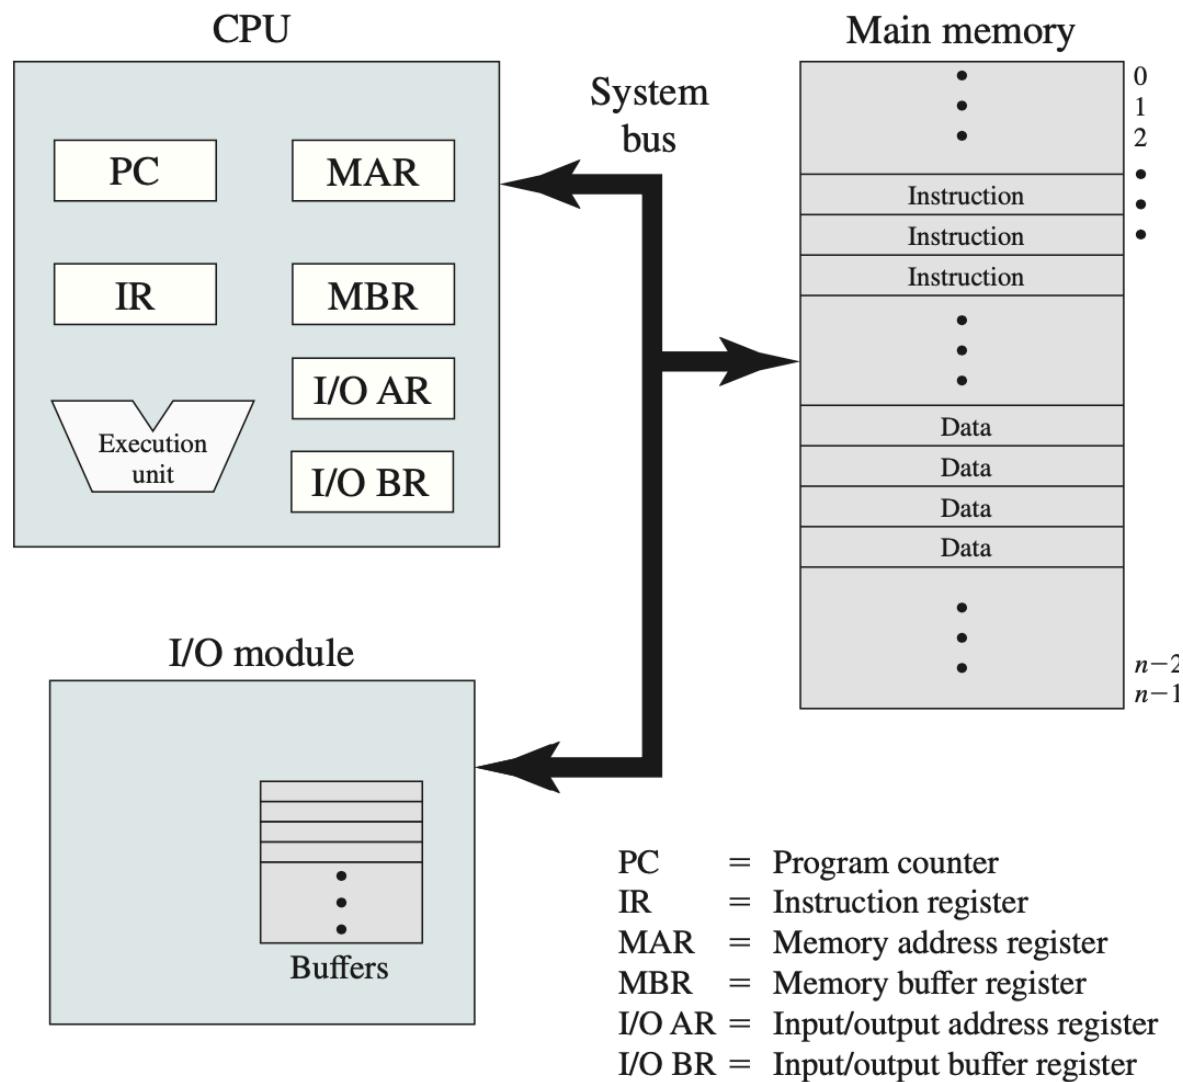


Figure 1.1 Computer Components: Top-Level View

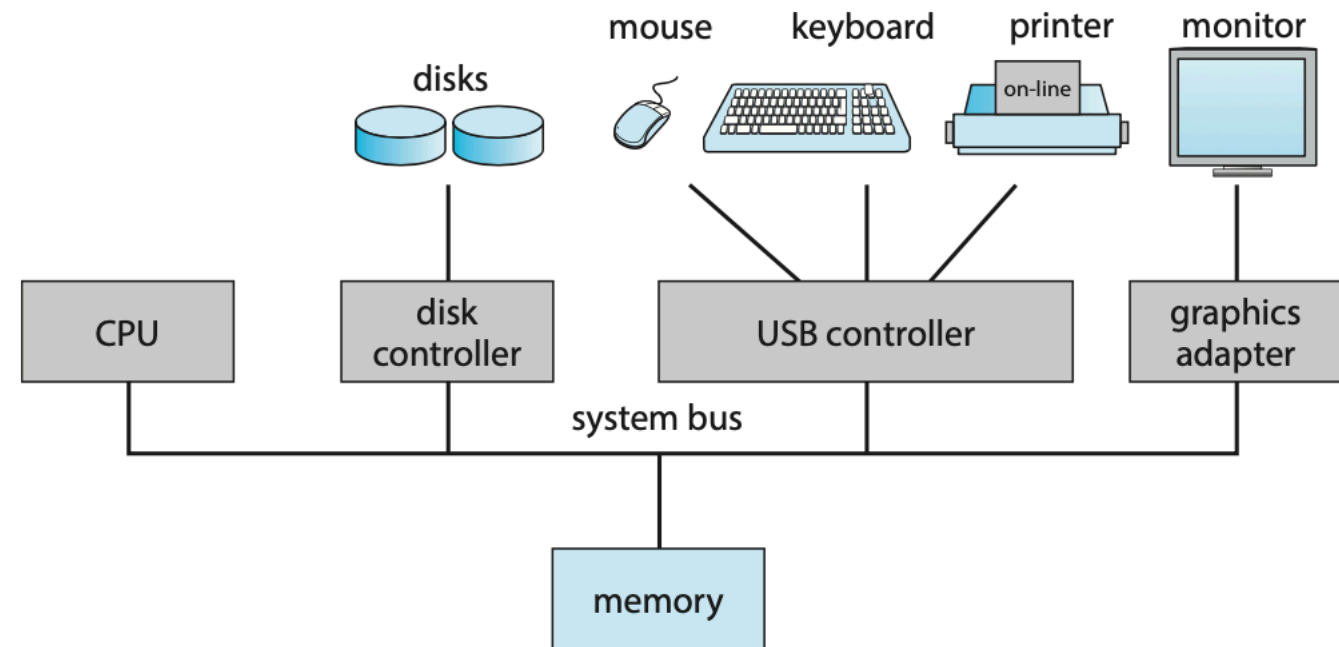



Figure 1.2 A typical PC computer system.



Storage Structure

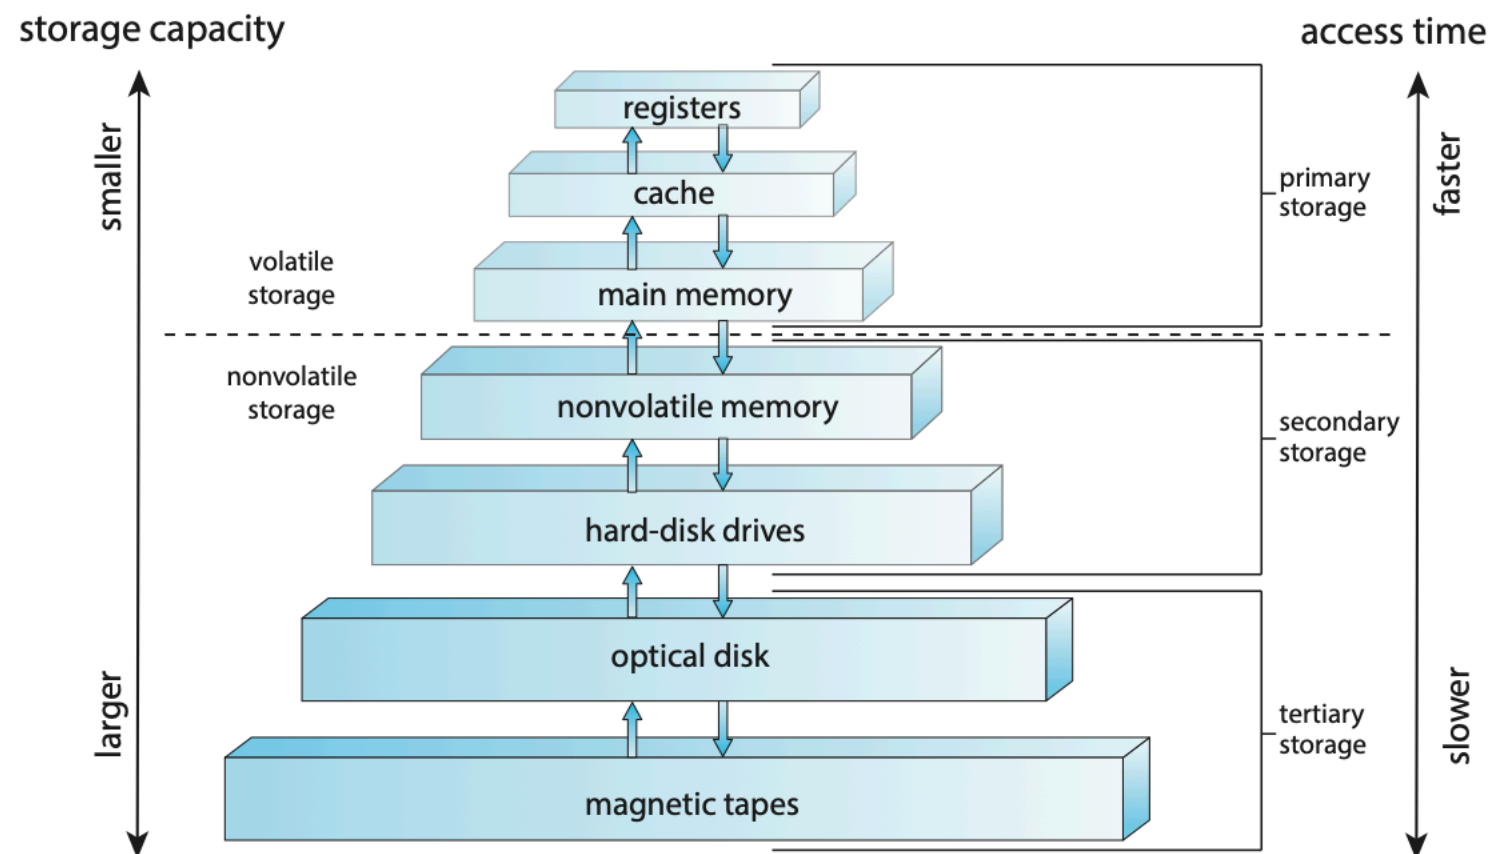


Figure 1.6 Storage-device hierarchy.

Computer Architecture

- Single-processor systems
 - Multiprocessor systems
 - Symmetric Multiprocessing (SMP)
 - Clusters
- **core** is the component that executes instructions and has registers for storing data locally




Figure 1.8 Symmetric multiprocessing architecture.




Figure 1.9 A dual-core design with two cores on the same chip.





Figure 1.11 General structure of a clustered system.

What happens when a program runs?



What is OS?

- User view: making sure the system operates correctly and efficiently in an **easy-to-use** manner.
- System view:
 - resource manager/allocator
 - control program
- OS as a **kernel** (the one program running at all times on the computer)

Three key ideas

- Virtualization
- Concurrency
- Persistence

Virtualization

Virtualization

- OS takes a physical resource (such as the processor, or memory, or a disk) and transforms it into a more general, powerful, and easy-to-use virtual form of itself
- provides some interfaces (APIs)
 - System calls (standard library to applications)

Virtualizing the CPU

cpu.c

```
#include <stdio.h>
#include <stdlib.h>
#include "common.h"

int main(int argc, char *argv[])
{
    if (argc != 2) {
        fprintf(stderr, "usage: cpu <string>\n");
        exit(1);
    }
    char *str = argv[1];

    while (1) {
        printf("%s\n", str);
        Spin(1);
    }
    return 0;
}
```

```
os@VM$ gcc -o cpu cpu.c -Wall
os@VM$ ./cpu
usage: cpu <string>
os@VM$ ./cpu testing
testing
testing
testing
testing
^C
```

Run multiple programs at once

```
os@VM$ ./cpu "A" & ./cpu "B" & ./cpu "C" & ./cpu "D" &
```

```
[1] 6596
```

```
[2] 6597
```

```
[3] 6598
```

```
[4] 6599
```

```
A
```

```
os@VM$ D
```

```
C
```

```
B
```

```
C
```

```
D
```

```
A
```

```
B
```

```
C
```

```
D
```

```
A
```

```
B
```

```
C
```

```
A
```

```
D
```

```
B
```

```
kill $(jobs -p)
```

```
[1] Terminated
```

```
[2] Terminated
```

```
[3]- Terminated
```

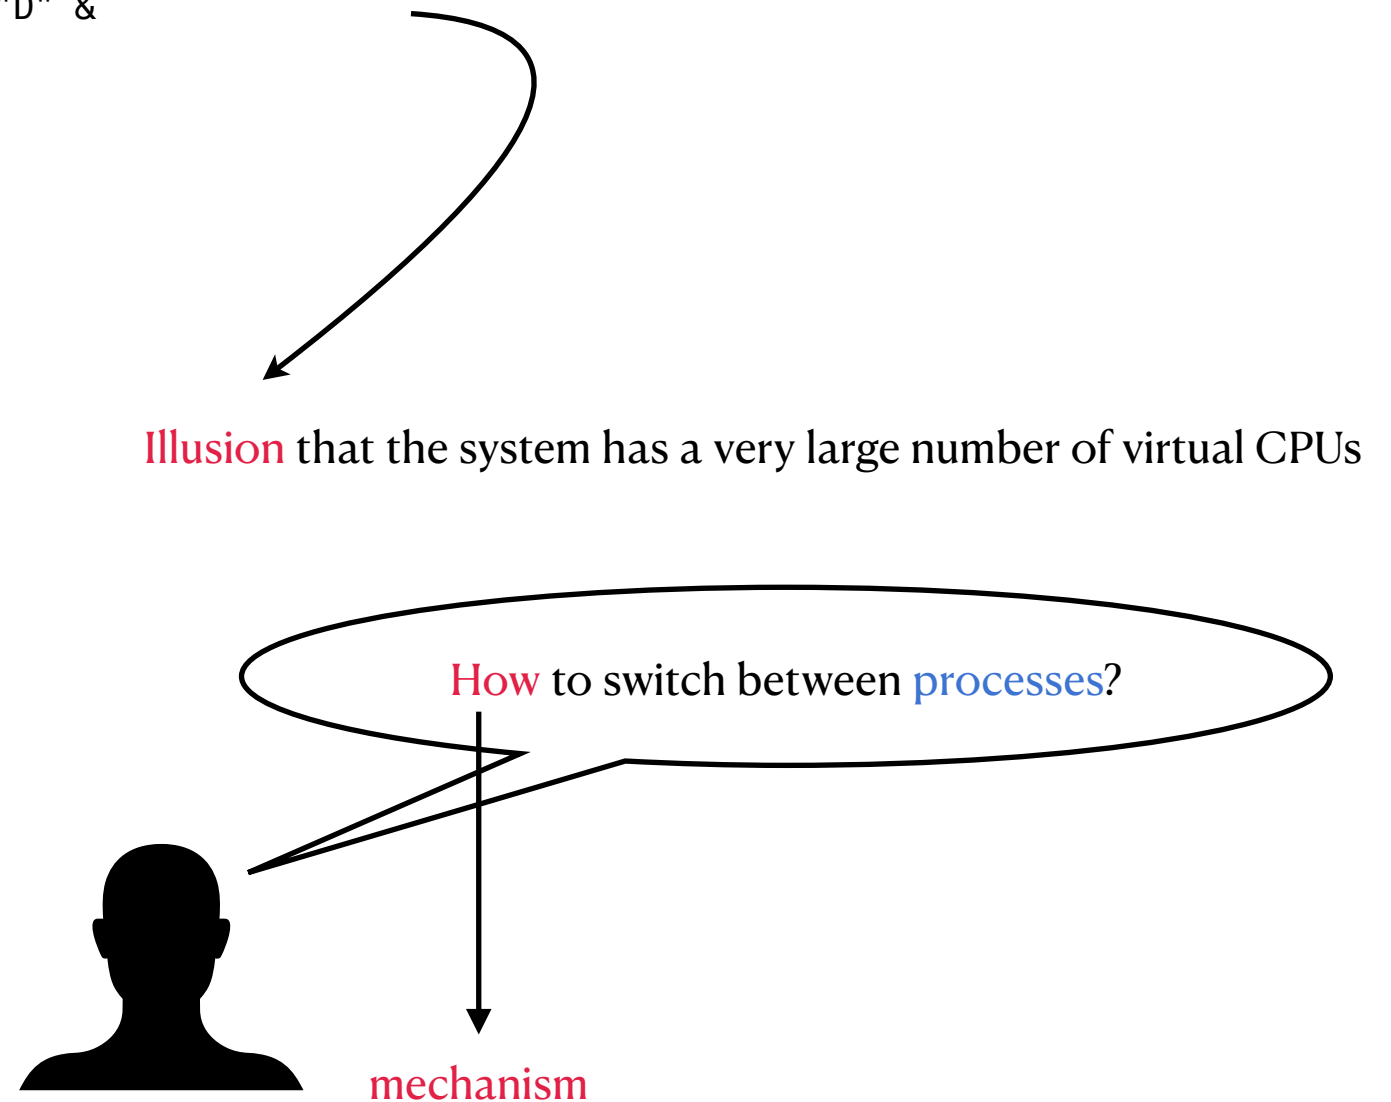
```
[4]+ Terminated
```

```
./cpu "A"
```

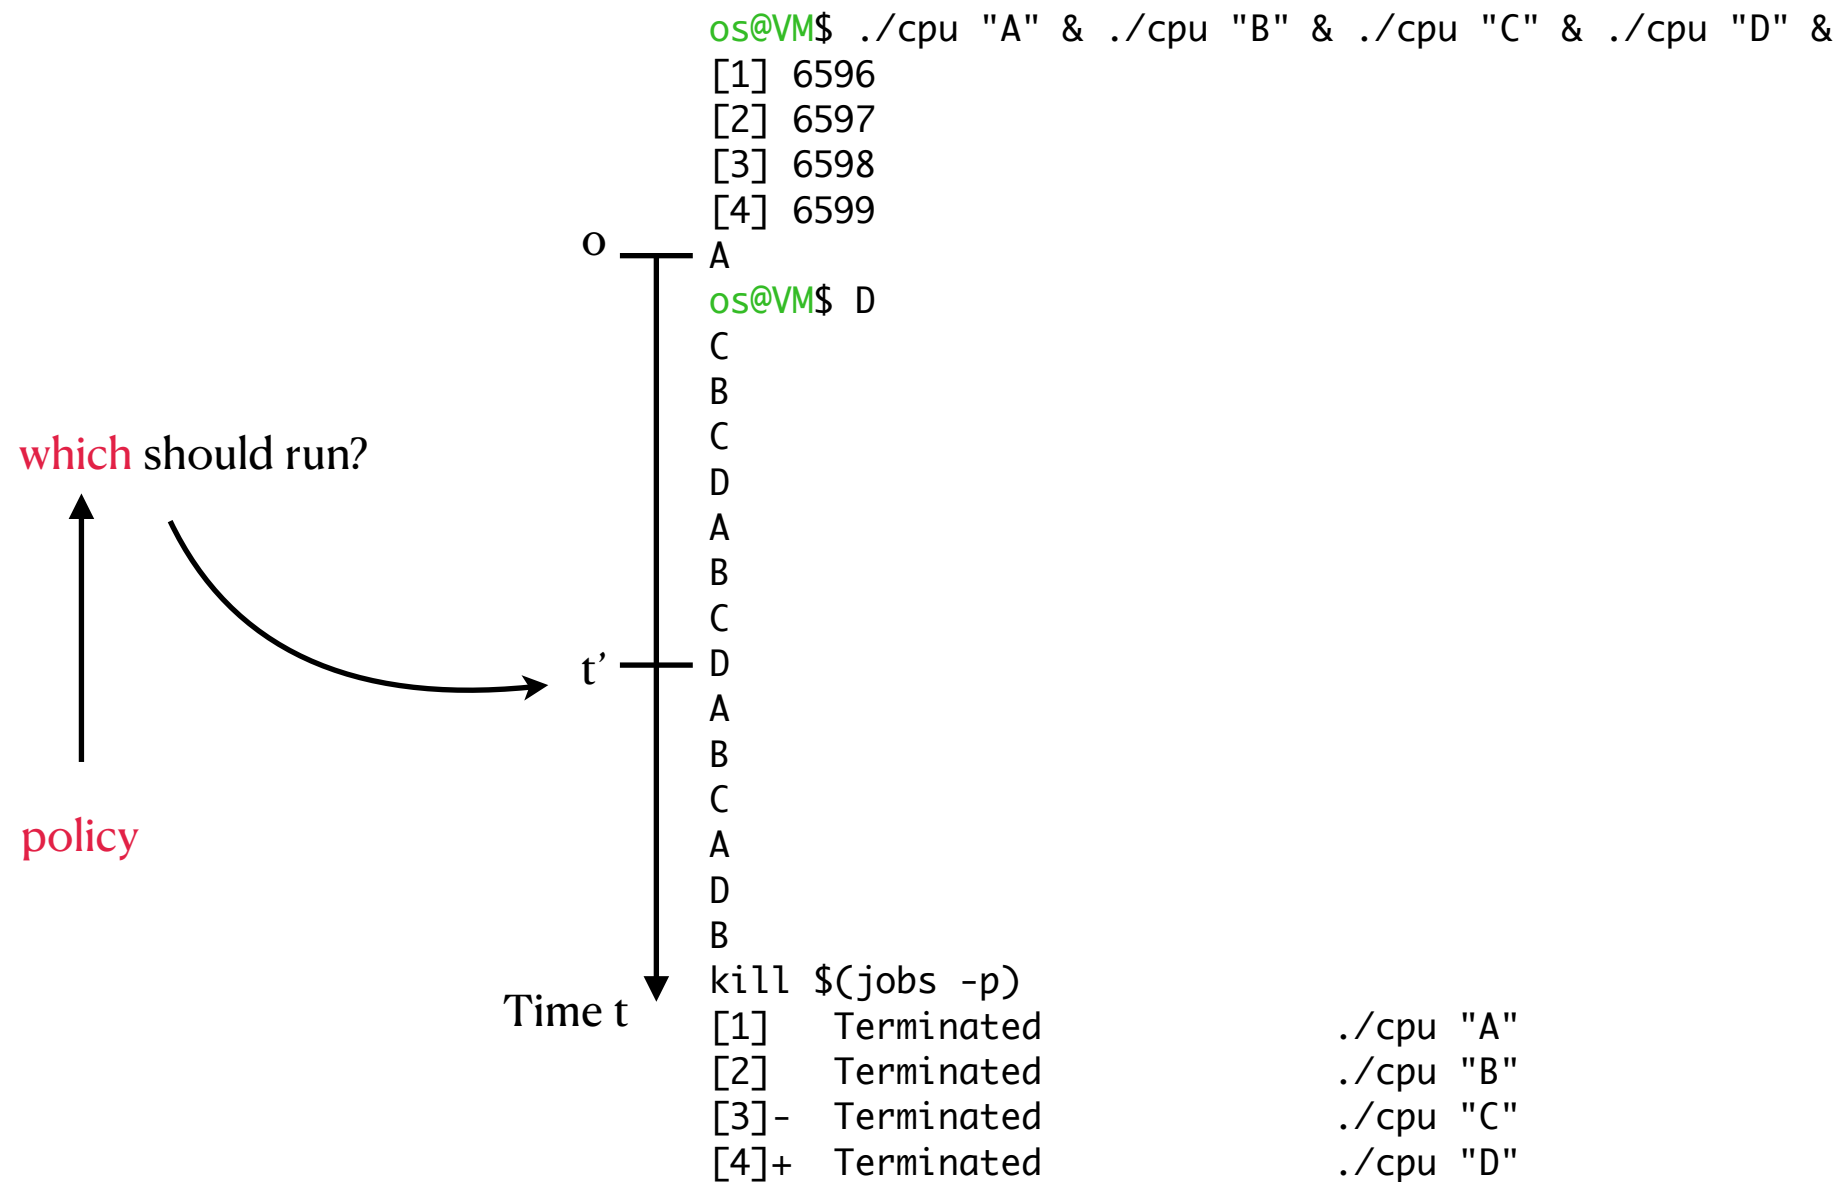
```
./cpu "B"
```

```
./cpu "C"
```

```
./cpu "D"
```



Run multiple programs at once



Virtualizing Memory

mem.c

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include "common.h"

int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, "usage: mem <value>\n");
        exit(1);
    }
    int *p;
    p = malloc(sizeof(int));
    assert(p != NULL);
    printf("(pid) addr pointed to by p: %p\n", (int) getpid(), p);
    *p = atoi(argv[1]); // assign value to addr stored in p
    while (1) {
        Spin(1);
        *p = *p + 1;
        printf("(pid) value of p: %d\n", getpid(), *p);
    }
    return 0;
}
```

```
os@VM$ gcc -o mem mem.c -Wall
```

```
os@VM$ ./mem 5
```

```
(6985) addr pointed to by p: 0x55a3492fa2a0
```

```
(6985) value of p: 6
```

```
(6985) value of p: 7
```

```
(6985) value of p: 8
```

```
(6985) value of p: 9
```

```
(6985) value of p: 10
```

```
(6985) value of p: 11
```

```
^C
```

```
os@VM$ ./mem 0 & ./mem 0 &
```

```
[5] 3927
```

```
[6] 3928
```

```
(3927) addr pointed to by p: 0x5586b2a432a0
```

```
(3928) addr pointed to by p: 0x5653bb05d2a0
```

```
(3927) value of p: 1
```

```
(3928) value of p: 1
```

```
(3927) value of p: 2
```

```
(3928) value of p: 2
```

```
(3927) value of p: 3
```

```
(3928) value of p: 3
```

```
kill $(jobs -p)
```

Virtualization Summary

- OS as a resource manager
 - Share CPU —> virtualizing CPU
 - Share memory —> virtualizing memory

Concurrency

Concurrency

- Problems arise when working on many things at once (i.e., concurrently)
- Multi-threaded programs exhibit the same problems

threads.c

```
#include <stdio.h>
#include <stdlib.h>
#include "common.h"
#include "common_threads.h"

volatile int counter = 0;
int loops;

void *worker(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        counter++;
    }
    return NULL;
}

int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, "usage: threads <loops>\n");
        exit(1);
    }
    loops = atoi(argv[1]);
    pthread_t p1, p2;
    printf("Initial value : %d\n", counter);
    Pthread_create(&p1, NULL, worker, NULL);
    Pthread_create(&p2, NULL, worker, NULL);
    Pthread_join(p1, NULL);
    Pthread_join(p2, NULL);
    printf("Final value : %d\n", counter);
    return 0;
}
```

```
os@VM$ gcc -o 2thread threads.c -Wall -pthread
os@VM$ ./thread 1000
Initial value : 0
Final value : 1000
os@VM$ ./1thread 1000
Initial value : 0
Final value : 1000
os@VM$ ./2thread 1000
Initial value : 0
Final value : 2000
os@VM$ ./2thread 10000
Initial value : 0
Final value : 20000
os@VM$ ./2thread 100000
Initial value : 0
Final value : 193888
```

Persistence

Persistence

- No virtualized disk for each application
- Handling problems with **journaling** or **copy-on-write**
- Topics: I/O devices, disks, RAIDs, file systems

io.c

```
#include <stdio.h>
#include <unistd.h>
#include <assert.h>
#include <fcntl.h>
#include <sys/stat.h>
#include <sys/types.h>
#include <string.h>

int main(int argc, char *argv[]) {
    int fd = open("/tmp/file", O_WRONLY | O_CREAT | O_TRUNC, S_IRUSR | S_IWUSR);
    assert(fd >= 0);
    char buffer[20];
    sprintf(buffer, "hello world\n");
    int rc = write(fd, buffer, strlen(buffer));
    assert(rc == (strlen(buffer)));
    fsync(fd);
    close(fd);
    return 0;
}
```

```
os@VM$ gcc -o io io.c -Wall
```

```
os@VM$ cat /tmp/file
```

```
hello world
```


Design Goals

- Build up some abstraction
- **Performance**: minimize the overheads of the OS
- Provide **protection** between applications, & between OS and applications
- **Reliability**
- **Security**
- **Mobility**